

Credit Card Default Prediction & Analysis


GROUP 1

NUS

March 24, 2021

Motivation

Credit risk plays a major role in the **banking industry** business. However, with the growing number of credit card users, banks have been facing an escalating credit card default and delinquency rate.



Source: New York Fed, DBS

Motivation

Fortunately, large scale (behavior) data is naturally suitable for developing machine learning models. More economic researchers and data scientists are focusing on this area.

Research Papers:

- Forecasting and stress testing credit card default using dynamic models (2013)
- Credit Card Default Prediction using Machine Learning Techniques (2018)
- Enhanced Recurrent Neural Network For Combining Statistic and Dynamic Features for Credit Card Default Prediction (2019)
- ...

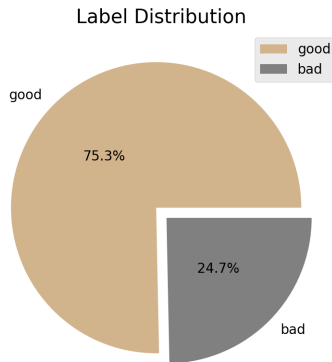
Data Science Competitions:

- Kaggle: Credit Card Default
- Kaggle: Default of Credit Card Clients
- "Magic Mirror Cup" Risk Management Algorithm Competition
- CMB Fintech Elite Training Camp (2020)
- **The 2nd Yipay Cup Big Data Modeling Competition**
- ...

Dataset

From "The 2nd Yipay Cup Big Data Modeling Competition"

- Supervised learning
- Imbalanced classification problem
- Real but desensitized data
- 47782 individuals
 - Train and Validation set: 80%
 - Test set 20%
- 3 data frames:
 - base
 - transaction
 - operation



Dataset: base.csv

DataFrame 1: base.csv → Static information

| | user | sex | age | provider | level | ... | product4_amount | product5_amount | product6_amount | product7_cnt | product7_fail_cnt |
|---|-------------|------------|-------|------------|------------|-----|-----------------|-----------------|-----------------|--------------|-------------------|
| 0 | Train_12996 | category 0 | 24889 | category 0 | category 2 | ... | level 0 | level 0 | level 1 | 24712 | 24712 |
| 1 | Train_39734 | category 0 | 24859 | category 0 | category 2 | ... | level 0 | level 0 | level 14 | 24712 | 24706 |
| 2 | Train_09006 | category 0 | 24931 | category 0 | category 2 | ... | level 0 | level 0 | level 20 | 24712 | 24706 |
| 3 | Train_35097 | category 1 | 24938 | category 2 | category 2 | ... | level 0 | level 0 | level 1 | 24712 | 24712 |
| 4 | Train_47615 | category 0 | 24853 | category 0 | category 1 | ... | level 0 | level 0 | level 1 | 24706 | 24706 |

5 rows × 46 columns


Encoding:

- nominal variables few values (e.g. sex): LabelEncoder
- nominal variables many values (e.g. city, province): CountEncoder
- ordinal or numerical variables: Transfer into integer

Feature Engineering:

- second order addition of numerical variables
- ratios from practical experience or economic intuitive


Dataset: base.csv



(a)



(b)



(c)

Figure: Box plots of generated feature

Now, We get 233 features.

Dataset: trans.csv

DataFrame 2 & 3: trans.csv & op.csv → Dynamic information

- Time series data starts from some time point
- Each user have different length of behavior

| | user | platform | tunnel_in | tunnel_out | ... | ip | type2 | ip_3 | tm_diff |
|---|-------------|------------------|------------------|------------------|-----|------------------|------------------|------------------|----------------------------|
| 0 | Train_13770 | 46c69cbbce5f1568 | b2e7fa260df4998d | 6ee790756007e69a | ... | NaN | 11a213398ee0c623 | NaN | 19 days 09:02:45.000000000 |
| 1 | Train_13770 | 46c69cbbce5f1568 | b2e7fa260df4998d | 6ee790756007e69a | ... | NaN | 11a213398ee0c623 | NaN | 19 days 09:03:58.000000000 |
| 2 | Train_08351 | 46c69cbbce5f1568 | b2e7fa260df4998d | 6ee790756007e69a | ... | f10a09fe9e522a47 | 11a213398ee0c623 | ee386d6f9fe45d0d | 18 days 11:06:49.000000000 |
| 3 | Train_08351 | 42573d7287a8c9c2 | NaN | 6ee790756007e69a | ... | NaN | NaN | NaN | 26 days 09:52:51.000000000 |
| 4 | Train_08351 | 42573d7287a8c9c2 | NaN | 6ee790756007e69a | ... | NaN | NaN | NaN | 26 days 07:50:05.000000000 |

5 rows × 10 columns

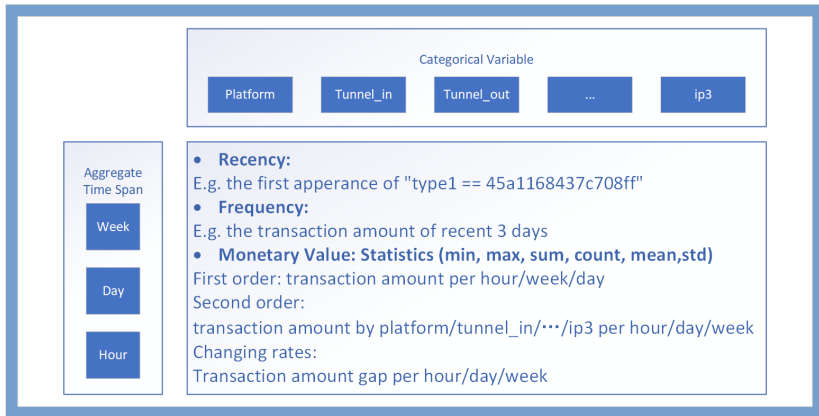
Main Issue

Represent the vector of each user's behavior record ($N \times 10$) into an aggregation vector ($1 \times M$), then we can concatenate them with features from base.csv, and feed into model.

Dataset: trans.csv

Feature Engineering:

Tracking user's **Recency**, **Frequency**, **Monetary Value**



The diagram illustrates feature engineering options for the 'trans.csv' dataset. It is organized into three main sections within a light blue border:

- Categorical Variable** (top right): A horizontal row of five blue buttons labeled 'Platform', 'Tunnel_in', 'Tunnel_out', '...', and 'ip3'.
- Aggregate Time Span** (middle left): A vertical stack of three blue buttons labeled 'Week', 'Day', and 'Hour'.
- Feature Descriptions** (bottom right): A light blue box containing three bullet points:
 - Recency:** E.g. the first appearance of "type1 == 45a1168437c708ff"
 - Frequency:** E.g. the transaction amount of recent 3 days
 - Monetary Value: Statistics (min, max, sum, count, mean, std)**
 - First order: transaction amount per hour/week/day
 - Second order: transaction amount by platform/tunnel_in/.../ip3 per hour/day/week
 - Changing rates: Transaction amount gap per hour/day/week

Dataset: trans.csv

Feature Engineering: (mainly focusing on amount)

- Manually written functions

```
# define a function to calculate users' transaction amount linked with feature 'hours'
def gen_user_window_amount_features(df, window):
    group_df = df[df['hour'] > window].groupby('user')['amount'].agg([
        ('user_amount_mean_{}'.format(window), 'mean'),
        ('user_amount_std_{}'.format(window), 'std'),
        ('user_amount_max_{}'.format(window), 'max'),
        ('user_amount_min_{}'.format(window), 'min'),
        ('user_amount_sum_{}'.format(window), 'sum'),
        ('user_amount_med_{}'.format(window), 'median'),
        ('user_amount_cnt_{}'.format(window), 'count')
    ]).reset_index()
    return group_df
```

- Word Embedding: word2vec

Now, we get 671 features.

Feature Engineering: (no amount here)

- Manually written functions
- Word Embedding: word2vec
- Truncated svd on TF-IDF (`n_components=10`)
- Truncated svd on CountVectorizer (`n_components=10`)

Now, we get 994 features.

Baseline Model: LightGBM (5-fold CV)

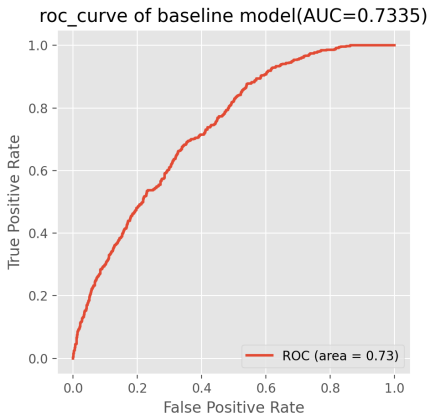


Figure: ROC curve of baseline model




Figure: Feature importance

Working On

- Model Tuning
- Try other models, e.g. XGBoost
- Stacking
- Oversampling (to handle imbalanced learning issue), e.g. SMOTE